

Classifying Breast Cancer Images Using Machine Learning Methods

Yuxuan Chen | Yuan Meng | Paula Wu

Introduction & Data preprocessing

Breast cancer is one of the most common invasive cancers in women in the United States, second only to skin cancer [1]. Starting from different parts of the breast, breast cancer is usually marked by lumpiness or swells in the breasts and surrounding tissues. However, it is important to distinguish benign breast tumors from malignant ones, as non-cancer breast tumors are usually abnormal growths that do not spread outside of the breast. They are not life-threatening, even though some types of benign breast lumps can increase a women's risk of getting breast cancer [2]. On the other hand, malignant tumors are aggressive and deadly. Fortunately, the prognosis of the disease has been greatly improved once it is detected. Therefore, it's important to have diseased tissue accurately diagnosed, as misdiagnosis may lead to delayed intervention or may cause unnecessary stress on a patient.

Data Overview

The original dataset contains 569 observations and 33 columns [3]. Among the 569 observations, there are 212 malignant observations and 357 benign observations. Thirty out of the 33 columns are predictors regarding imaging features, computed from a digitized image of a fine needle aspirate (FNA) of a breast mass [4]. These features describe the characteristics of the cell nuclei present in the image. Ten real-valued features are computed for each cell nucleus:

- radius: mean of distances from the center to points on the perimeter
- texture: standard deviation of gray-scale values
- perimeter: perimeter of the tumor image
- area
- smoothness: local variation in radius lengths
- compactness: computed as $\frac{\text{perimeter}^2}{\text{area}} - 1.0$
- concavity: severity of concave portions of the contour
- concave points: number of concave portions of the contour
- symmetry
- fractal dimension: computed as $\frac{\text{coastline approximation}}{\text{perimeter}} - 1$

For each variable, its mean, standard error, and “worst” or largest (mean of the three largest values) of these features computed for each image are included in this dataset. In this project, our goal is trying to determine which features contribute to a better diagnosis of breast tissues images.

Data preprocessing

By the time we got the dataset, it is quite clean and complete. Therefore major data preprocessing of this dataset only entails factorizing outcome variables (the benign class as control) and removing the ID column.

Exploratory Data Analysis

Correlation Plots

After data preprocessing, we plot the correlation plot (**Figure 1**) of all the predictors. We can tell from the graph that there is a high correlation across the predictors. This is because our predictors include mean, standard deviation, and the largest values of the distributions of 10 features, which means that some predictors can be calculated from others. The plot shows that the size of the core tumor has a strong correlation with tumor perimeter, radius, and area in the breast. It also shows tumor compactness is most correlated with its concavity and symmetry. Though the dataset has high multicollinearity, we choose not to trim any predictor given that we don't know much about the mechanisms of image classification.

Feature Plots

According to the density featurePlot visualization of each feature's mean (**Figure 2**), we can see that the distributions of benign versus malignant tumors are not that overlapping, suggesting a relatively clear distinction within the data. More specifically, patients with malignant tumors tend to have higher mean values on "compactness", "radius", "concavity", "perimeter", "area", and "concave points" of the tumor images.

K-mean

We use K-mean clustering to partition the observation into 2 clusters. By utilizing the `fviz_nbclust()` function, we can visualize and determine the optimal number of clusters equals to 2 (**Figure 3**), which corresponds to the number of categories of the response variable (M = malignant vs. B = benign). The **K-mean** plot shows the distribution of clusters on two principal components - PC1 and PC2. With each response labeled by its true category, we can see that the malignant cases are mainly in cluster 1, while the benign cases are mainly in cluster 2.

Models

All 30 continuous predictors are included to predict the outcome variable - "diagnosis". A 70:30 test-train split was applied and yielded 399 rows of training data and 170 rows of testing data. The 30 predictors are the "mean", "standard error", and "worst" or largest of 10 real-valued features. A total of 8 models, including penalized logistic regression, KNN, MARS, linear discriminant analysis, CART, random forest, boosting, and support vector machine (both linear and radial kernels), were used to fit the data. Data are centered and scaled during model training using `caret`, if necessary. Considering the computation cost, we use 10-fold cross-validation to find the optimal values for models involving tuning parameters. The process to select tuning parameters is shown in **Figure 4**. Last but not the least, we use resampling method to compare the fitted models and use an unseen test dataset to examine the model fit.

Penalized Logistic regression

Due to the nature of the binary outcome of the dataset, penalized logistic regression can be fit to assess the performance of linear decision boundaries. All 30 predictors are included to fit the model. The best-tuned model has $\alpha = 0.45$ and $\lambda = 0.004$ as hyperparameters.

KNN

KNN is a non-linear relation model to predict class labels among k neighbors. It is also a black-box model that is easy to implement and has good classification performance. One disadvantage of KNN is that if we wrongly choose the optimal value of k , the model can be under or overfitted to the data. The best tuning parameter for KNN is when the number of neighbors equals 28.

MARS

MARS model can automatically model non-linearities and interactions between variables. This model is well-fit for high-dimensional problems. There are two tuning parameters associated with this model: degree of interaction and number of retained terms. So, we performed a grid search and cross-validation to identify the optimal combination of these hyper-parameters. The combination of the best tuning parameters is when there are 8 retained terms (including intercept) and 1 degree of interaction.

Linear Discriminant Analysis (LDA)

As its name suggests, LDA is a linear classification method. This model projects a feature space onto a smaller subspace and classifies data points based on the nearest centroid; therefore does not involve any hyperparameters. It also assumes that the features follow Gaussian distribution [5]. While one advantage of LDA is its robustness to data that has two or more response classes, it is not the case in this breast cancer study that has binary outcomes.

Naive Bayes

With mechanisms similar to LDA, Naive Bayes becomes handy when there is a large number of predictors. However, this model does require feature independence, an assumption that can be easily violated here - as shown in EDA. Therefore, this method will only be briefly discussed in this section, and no model is fitted.

Classification and Regression Tree (CART)

CART is a tree-based method that grows branches in a top-down greedy fashion through recursive binary splitting and aims to partition the predictor space into several axis-parallel regions. When building the model, we chose to use the Gini index, a measure of node purity, as the stopping criteria for splitting. We tune the hyperparameter `Cp` through cross-validation and find that the optimal `Cp` equals 0.0072. As a single tree, it does have some drawbacks such as instability and high variance. Nevertheless, this model is included as a baseline for comparison with more complex tree-based models presented next.

Random Forest

As an ensemble model, the random forest is flexible and has a better chance to capture a complex underlying true model by adopting the “wisdom of crowds” idea [6]. The model contains multiple trees that are built on bootstrapped random observations, and for each split within each tree, a fresh subset of predictors will be selected. Therefore, the random forest can decorrelate the trees and thus reduce the variance compared to a single tree. Similar to CART, we adopt the Gini index to control splitting. There are 2 hyperparameters to tune in the random forest model: `mtry` and `minimum node size`. According to the `bestTune` object of the result, the model achieved its highest AUC when `mtry` = 1 and `min.node.size` = 2, indicating a small tree for each tree built.

Boosting

Boosting is a similar approach to the random forest, except that the trees are grown sequentially - each tree is grown using information from previously grown trees. Three hyperparameters, namely the total number of trees B , numbers of splits d , and shrinkage λ , are tuned. We use grid search, together with cross-validation, to find the optimal value for each tuning parameter. The optimal model that has the largest AUC is achieved when there are 4 degrees of interaction, 5000 trees fit, and 0.002 shrinkage.

Support Vector Machine (SVM with Linear and Radial Kernels)

SVM works by mapping data to a high-dimensional feature space so that data points can be categorized by the separating hyperplanes. The separating hyperplane that makes the biggest gap or margin between the two classes is selected. The decision function is fully specified by the support vectors (SVs), data points that

lie closest to the decision surface, and the elements in the training set that would change the position of the hyperplane if removed [7]. For the linear kernel, support vector classifiers with one tuning parameter `cost` (tuned to be 0.424) are used to build a linear boundary. Since most real data sets will not be fully separable by a linear boundary, the radial kernel is used. At its core, the radial kernel indicates that support vector machines can construct classification boundaries that are nonlinear in shape. When tuning for the radial kernel, we add another tuning parameter - `sigma`. After cross-validation, the optimal `cost` is found to be 1.881 while `sigma` is 0.045.

Results

Model Comparison: Resampling & Model Performance

After fitting the training dataset to all 8 models, we measure their resampling AUC scores. The mean resampling AUC scores, sensitivity, and specificity are included in **Figure 5**. In general, all the models we fit have quite a high resampling AUC scores, and the top 4 models (radial SVM, penalized logistic regression, linear SVM, and random forest) are extremely close in their mean AUC values. On the other hand, the CART model has the lowest mean AUC with higher variance compared to other models, which verifies that as a single tree model, CART is unstable.

Interestingly, although all 8 models have relatively high AUC, their sensitivity and specificity are more widespread. Among the top four models with the highest resampling AUC scores, SVM with radial kernel has the lowest mean sensitivity and highest mean specificity.

It is worth noting that neither the training nor the resampling AUCs should be adopted as the ultimate criterion when selecting the models. A model with large training AUC does not mean it has the best fit. Instead, we should re-run the models with unseen testing data to see if the model overfits the training data. The test AUCs are examined with results presented in **Table 1**. There isn't much difference in both ranking and values of test AUC and train/resampling AUCs, which suggests a good fit.

Although the radial SVM model has the largest AUC, its mean sensitivity is below most of our other models, which can potentially lead to the delayed intervention of malignant breast tumors. Therefore, considering all facts above, we choose the penalized logistic regression model as our optimal final model, with high sensitivity, specificity, and high interpretability.

Variable Importance

We select the top 4 models with the largest AUC scores to display their variable importance. We can see that the radial SVM model, linear SVM model, and random forest model all identify the "worst" perimeter (of the tumor), the "worst" radius (mean of distances from the center to points on the perimeter), the "worst" area, the "worst" concave points (number of concave portions of the contour), the mean of the concave points (number of concave portions of the contour), and the mean of the perimeter (of the tumor) as their top 6 important variables. And penalized logistic regression model identifies the variables of the standard error of fractal dimension ("coastline approximation" - 1), the standard error of smoothness (local variation in radius lengths), the "worst" smoothness (local variation in radius lengths), the mean of the concave points (number of concave portions of the contour), the standard error of compactness ($\text{perimeter}^2 / \text{area} - 1.0$), and the "worst" concave points (number of concave portions of the contour) as its top 6 important variables.

Conclusion & Discussion

As discussed above, we choose the Penalized Logistic Regression as our final model due to its relatively high performance and accessible interpretability. In the EDA section above, we stated that there is a high correlation across the predictors. By conducting penalized logistic regression, we can reduce the effect of high collinearity between predictors and thus make a better prediction.

The coefficients of the final model are shown in **Table 2**. Interpretation of the coefficient of **radius mean** is provided as follows: holding all other variables fixed, with every one unit of increase in the mean radius of the breast mass, the odds ratio of the imaging being categorized as malignant will increase by 5.65% (calculated as $(e^{0.055} - 1) \times 100$). While we don't provide the full interpretation of all coefficients, we do want to point out that coefficients of several variables - such as worst smoothness, SE of smoothness, and mean of concave points - are large in magnitude, indicating that the change in values of these predictors will result in a huge change in the odds ratio of classifying an image to be malignant. This coincides with what we observed in the variable importance section to some extent.

With respect to our goal mentioned in the previous section, we can conclude that our final model performs relatively well in distinguishing between malignant and benign tumors using imaging measures.

Last but not the least, we do notice that this is quite a well-separated data since all 8 trained models have achieved high model performance. Maybe in the future, we can apply the same methods to a different dataset to examine the model performance along with the important variables.

Appendix




Figure 1: Correlation Plot

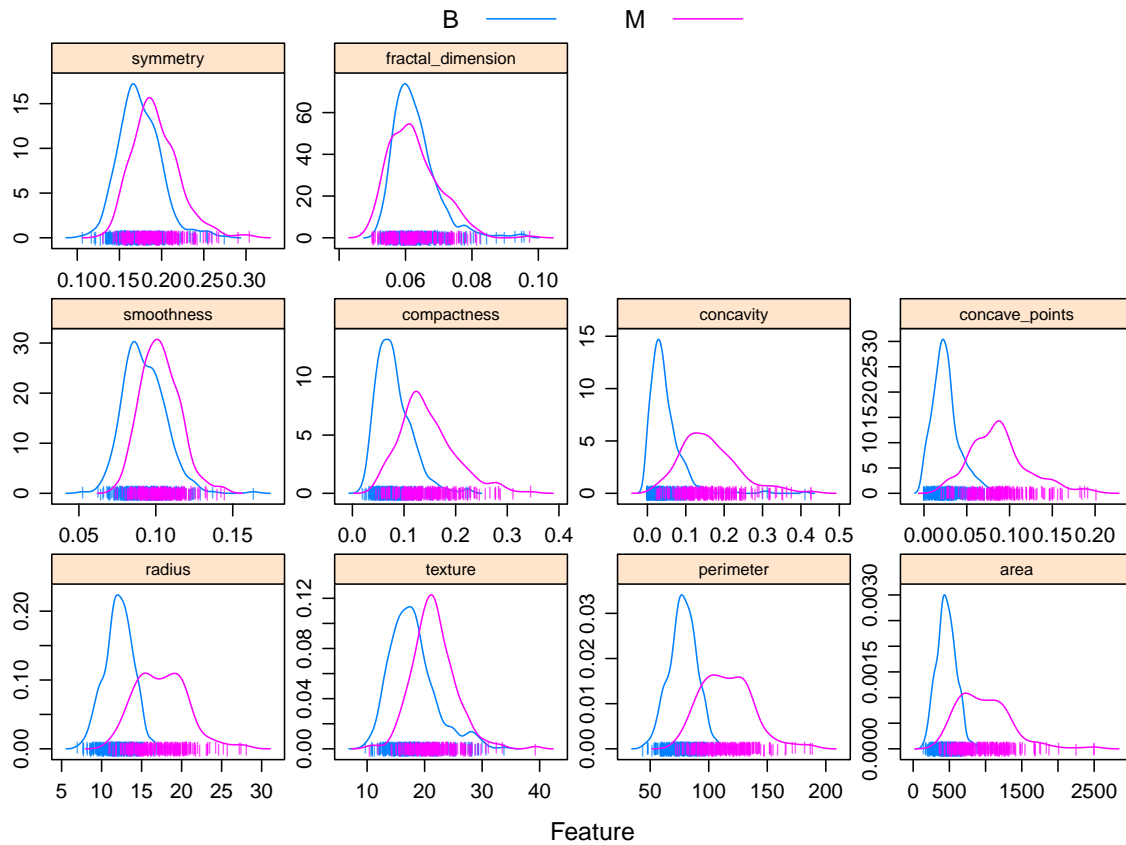


Figure 2: Feature Plots




Figure 3: K-mean Clustering

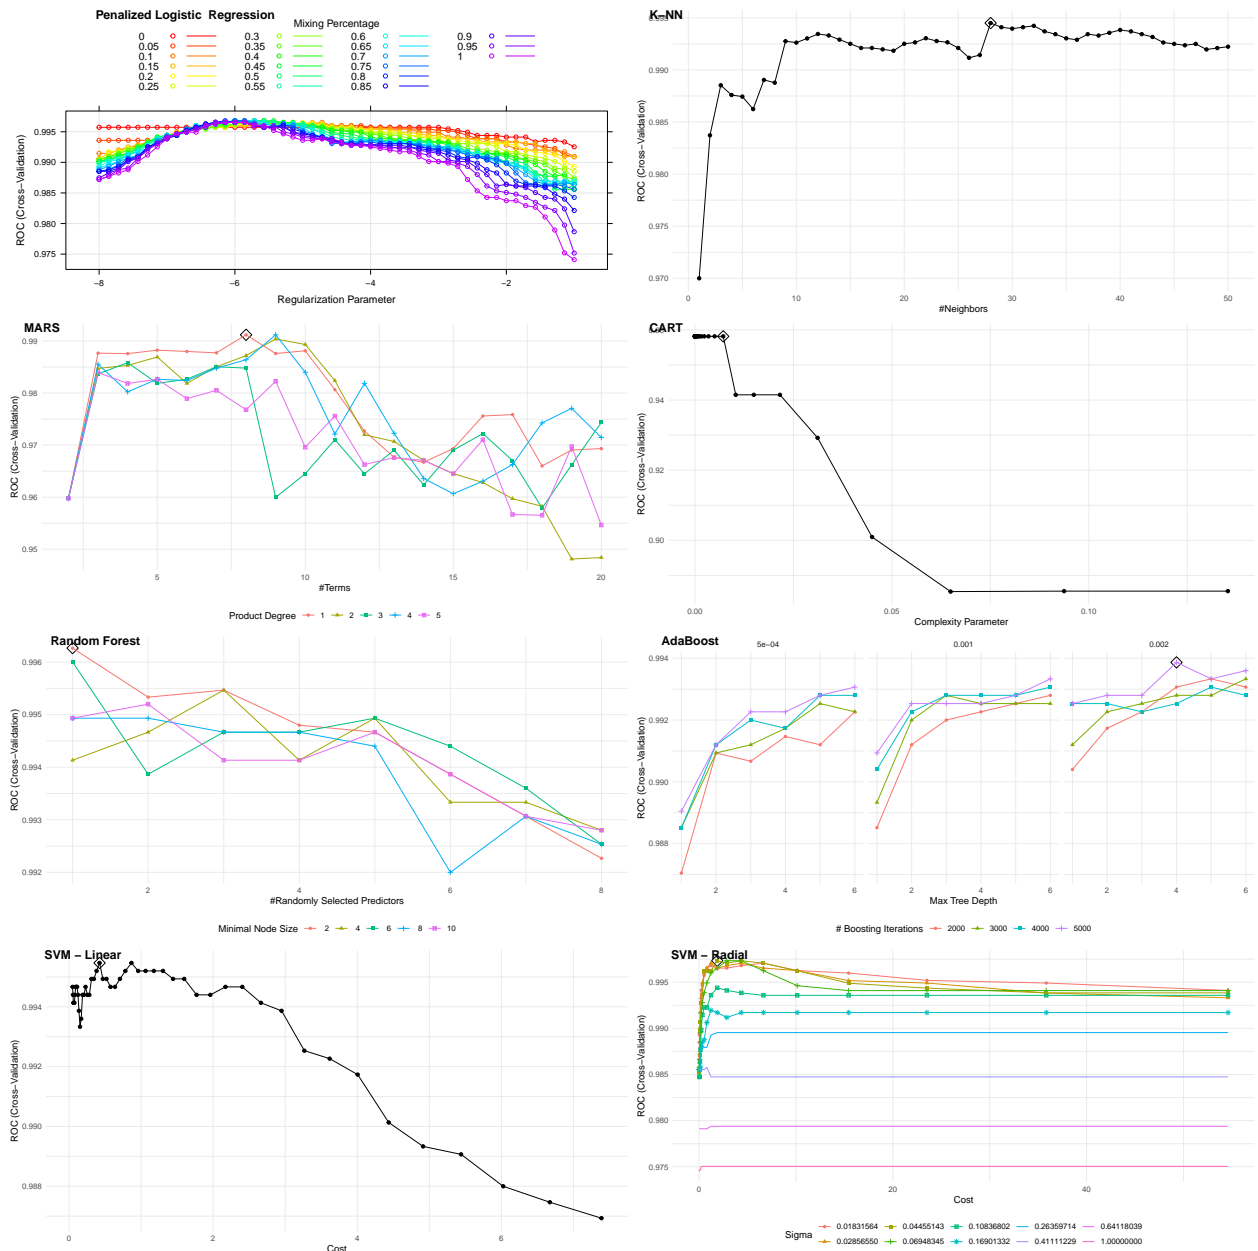


Figure 4: Model tuning




Figure 5: Resampling Results




Figure 6: Variable Importance

Table 1: testAUC scores

model	testAUC
SVML	0.9974781
Plr	0.9958463
LDA	0.9958463
MARS	0.9951046
SVMR	0.9949562
KNN	0.9913959
GB	0.9884290
RF	0.9872422
CART	0.9456312

Table 2: Coefficients of Penalized Logistic Regression Model

	coef
(Intercept)	-25.5059446
radius_mean	0.0557381
texture_mean	0.1050963
perimeter_mean	0.0072075
area_mean	0.0004757
concavity_mean	4.2732321
concave_points_mean	21.0406499
radius_se	3.9013407
perimeter_se	0.2671455
area_se	0.0103224
smoothness_se	-27.1800544
compactness_se	-17.2343871
symmetry_se	-9.6302843
fractal_dimension_se	-193.9228526
radius_worst	0.2049761
texture_worst	0.1229049
perimeter_worst	0.0245657
area_worst	0.0013846
smoothness_worst	27.1684565
concavity_worst	2.1546786
concave_points_worst	14.5823135
symmetry_worst	8.6031276
fractal_dimension_worst	1.0808845

Reference

- [1] Breast cancer basic information
- [2] What is breast cancer
- [3] [Data] (<https://www.kaggle.com/code/buddhiniw/breast-cancer-prediction/data>)
- [4] Breast Cancer Screening
- [5] LDA mechanism: Lecture note 9
- [6] Tree-based/ Ensemble models: Lecture note 12
- [7] SVM mechanisms: Lecture note 14
- [8] Actual implementation